



TOPIC G01

UNDERSTANDING
NOSQL DATABASE

UNDERSTAND NOSQL DATABASE

CONTENT

- Explain advantages and disadvantages of relational database
- Explain NoSQL and type of NoSQL databases
- Use MongoDB
- Explain NoSQL data modeling

Reference:

1. Dan Sullivan, Advanced NoSQL for Data Science (2017), LinkedIn Course
2. MongoDB documentation, <https://docs.mongodb.com/manual/core/data-modeling-introduction/>

RELATIONAL DATABASE - RECAP



Records are organized into tables.



Rows of tables are identified by keys.



Data spans multiple tables, which are linked by a join operation.



Transactions are ACID-compliant.



Data models are normalized.

Atomicity

- Each transaction is 'all or nothing'.
- If part of the transaction fails, then the entire transaction fails.
- The database state is left unchanged.

Consistency

- Any transaction will bring the database from one valid state to another.
- It satisfies all the rules and constraints that database.

Isolation

- The concurrent execution of transactions results in a system state that would be obtained if the transactions were executed serially.

Durability

- Once transaction has been committed, it will remain so.

RELATIONAL DATABASE

Advantages

- Comprehensive querying
- Normalized data
- Widely supported

Disadvantages

- Fixed schema
- Costly to join
- Limited data structures
- Difficult to scale

RELATIONAL DATABASE – MINIMIZE THE CONS

Denormalization

- Expand the number of columns in a single table

Partition

- Breaking up the database and store pieces of database in different servers

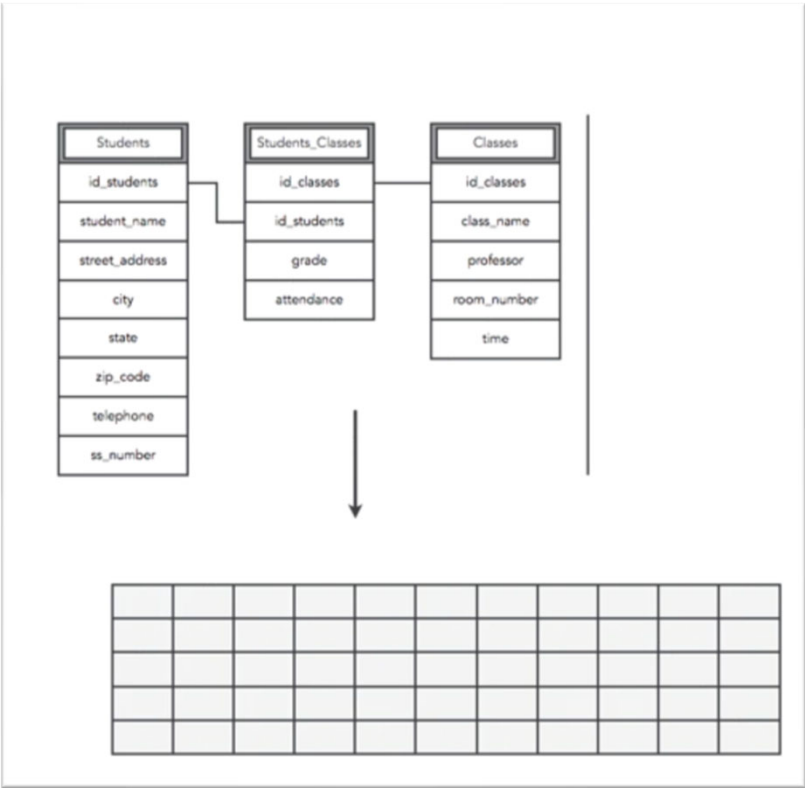
Replication

- Making copies of the database, store the copies in different servers

RELATIONAL DATABASE – MINIMIZE THE CONS

Denormalization

- Avoid joins
- Expand the number of columns
- Design table to include related data
- Query a single table
- **Improve read performance**



WHAT IS NOSQL DATABASE?

- Not using SQL
- Flexible Schema
 - No fixed-table definition
 - Fields are not standardized between records
- Support complex data structures
 - Nested values are common in NoSQL database
- Design with 'partition' in mind (Scalability)

NOSQL DATABASE TYPE

Key-value (eg Amazon S3)

- Based on key-value or dictionary data structures

Document (eg CouchDB, MongoDB)

- Multiple key-value pairs in a document

Wide Column (eg Cassandra, Google Bigtable)

- Organized in tables, rows and columns
- Columns are not fixed, can be different between rows and can be changed

Graph (Neo4j)

- Network of connected entities
- Entities are linked by edges

SHORT QUIZ

- Which type of NoSQL database is based on sets of nodes and edges between nodes?
(Choose: Document, Wide-Column, Graph)
- Which is not an advantage of NoSQL database?
(Choose: Support 'Joins', Support for large dataset, Support distributed database)



The database for modern applications

MongoDB is a general purpose, document-based, distributed database built for modern application developers and for the cloud era. No database makes you more productive.

 Try MongoDB free in the cloud!

[Get Started](#)



DOCUMENT DATA MODEL – DOCUMENT (~ RECORD)

Documents

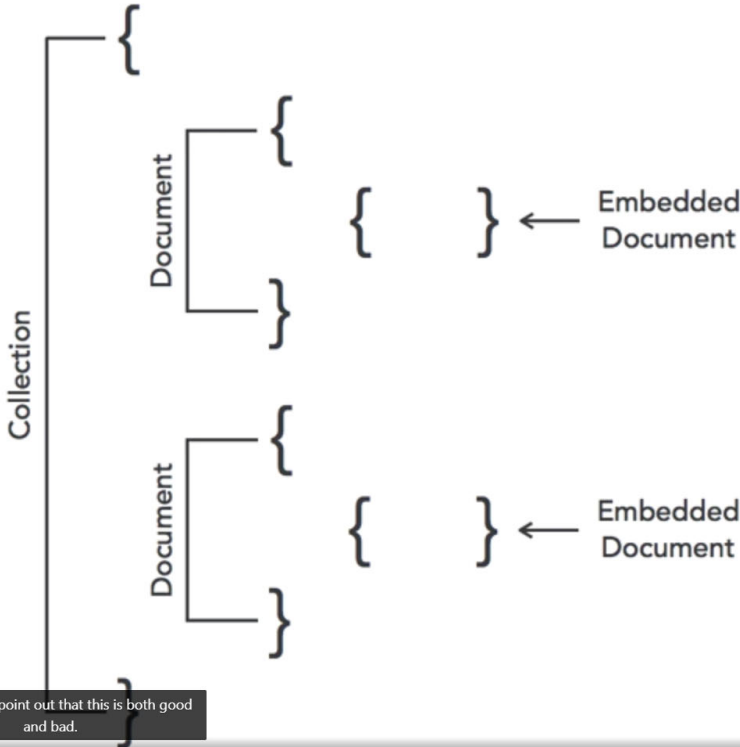
- Set of key-values
- Optional indexes
- Attributes vary across documents.
- Values may be complex structures.

```
{  
  orderID: 165345,  
  
  customer:{ FullName:'Morgan James',  
             Address: '321 First St',  
             City: 'Portland'},  
  
  orderItem:{ SKU: 5638126,  
              Desc: 'Small Pen'},  
  
  orderItem:{ SKU: 5639444,  
              Desc: 'Notepad'},  
  
  orderItem:{ SKU: 5627734,  
              Desc: 'Scissors'}  
}
```

DOCUMENT DATA MODEL – COLLECTION (~ TABLE)

Collections

- Set of 0 or more documents
- Each document has a unique ID.
- Schema inferred
- New attributes added by creating a document with attributes



RUN MONGODB COMMUNITY EDITION

Set up Mongo
Environment

- Create `\data\db`

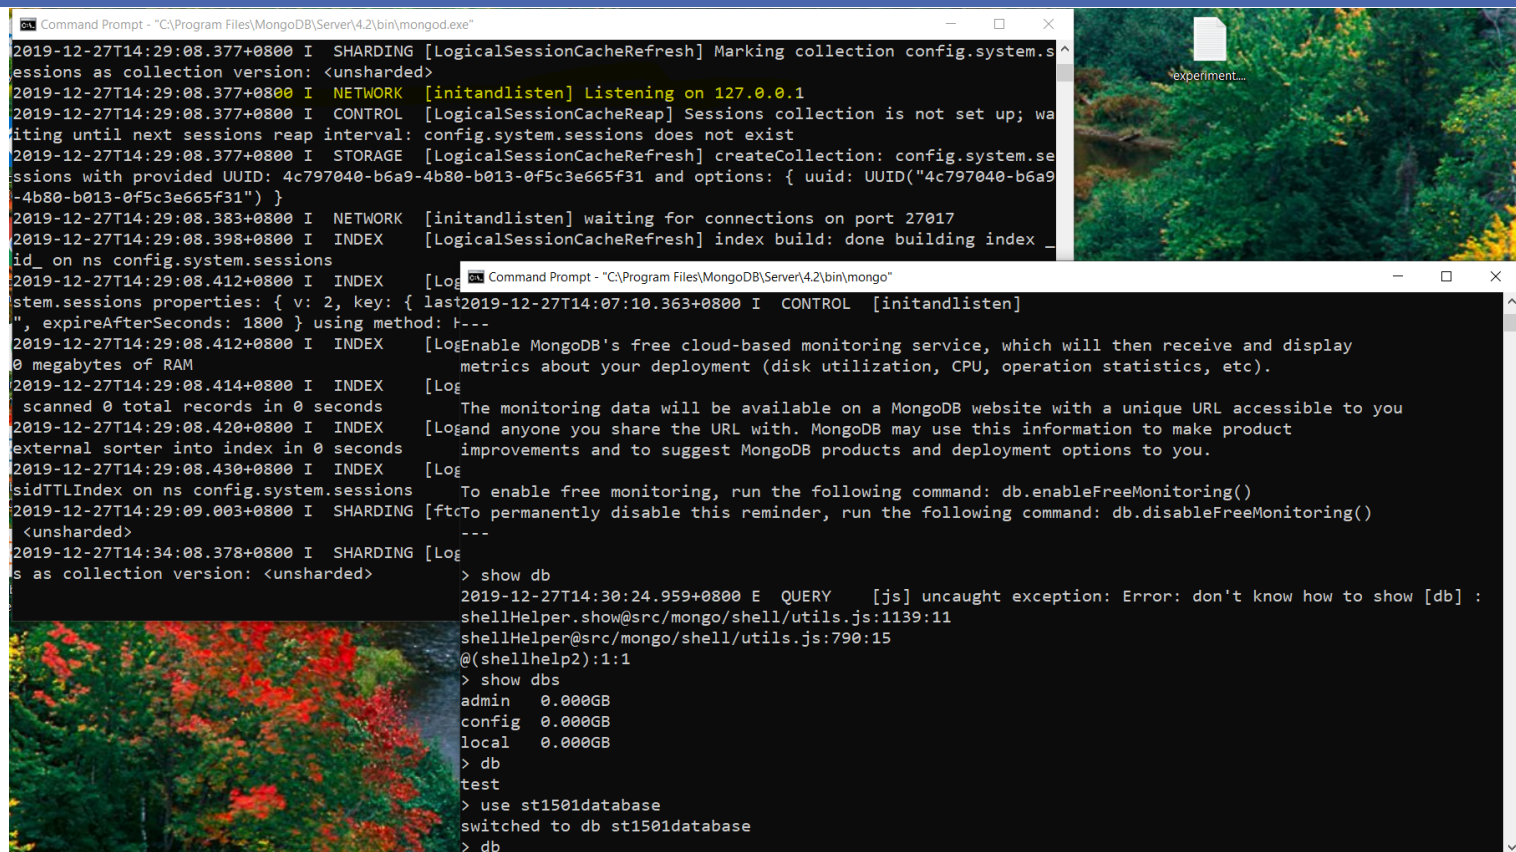
Start
MongoDB

- Run **mongod.exe**
- If not using default dbpath, specify `--dbpath actualpath`

Connect to
MongoDB via
Mongo Shell

- Run **mongo.exe**

RUN MONGODB COMMUNITY EDITION



```

Command Prompt - "C:\Program Files\MongoDB\Server\4.2\bin\mongod.exe"
2019-12-27T14:29:08.377+0800 I SHARDING [LogicalSessionCacheRefresh] Marking collection config.system.sessions as collection version: <unsharded>
2019-12-27T14:29:08.377+0800 I NETWORK [initandlisten] Listening on 127.0.0.1
2019-12-27T14:29:08.377+0800 I CONTROL [LogicalSessionCacheReap] Sessions collection is not set up; waiting until next sessions reap interval: config.system.sessions does not exist
2019-12-27T14:29:08.377+0800 I STORAGE [LogicalSessionCacheRefresh] createCollection: config.system.sessions with provided UUID: 4c797040-b6a9-4b80-b013-0f5c3e665f31 and options: { uuid: UUID("4c797040-b6a9-4b80-b013-0f5c3e665f31") }
2019-12-27T14:29:08.383+0800 I NETWORK [initandlisten] waiting for connections on port 27017
2019-12-27T14:29:08.398+0800 I INDEX [LogicalSessionCacheRefresh] index build: done building index _id_ on ns config.system.sessions
2019-12-27T14:29:08.412+0800 I INDEX [Log]
stem.sessions properties: { v: 2, key: { last2019-12-27T14:07:10.363+0800 I CONTROL [initandlisten]
", expireAfterSeconds: 1800 } using method: b-
2019-12-27T14:29:08.412+0800 I INDEX [Log]Enable MongoDB's free cloud-based monitoring service, which will then receive and display
metrics about your deployment (disk utilization, CPU, operation statistics, etc).
2019-12-27T14:29:08.414+0800 I INDEX [Log]
The monitoring data will be available on a MongoDB website with a unique URL accessible to you
2019-12-27T14:29:08.420+0800 I INDEX [Log]and anyone you share the URL with. MongoDB may use this information to make product
improvements and to suggest MongoDB products and deployment options to you.
2019-12-27T14:29:08.430+0800 I INDEX [Log]
To enable free monitoring, run the following command: db.enableFreeMonitoring()
2019-12-27T14:29:09.003+0800 I SHARDING [ftcTo permanently disable this reminder, run the following command: db.disableFreeMonitoring()
<unsharded>
2019-12-27T14:34:08.378+0800 I SHARDING [Log]
s as collection version: <unsharded>
> show db
2019-12-27T14:30:24.959+0800 E QUERY [js] uncaught exception: Error: don't know how to show [db] :
shellHelper.show@src/mongo/shell/utils.js:1139:11
shellHelper@src/mongo/shell/utils.js:790:15
@(shellhelp2):1:1
> show dbs
admin 0.000GB
config 0.000GB
local 0.000GB
> db
test
> use st1501database
switched to db st1501database
> db

```

MONGODB: DATABASE & COLLECTION

- MongoDB stores JSON **documents** in **collection**, collections in **database**.

```
{
  name: "a1",
  age: 18,
  status: "D",
  groups: [ "politics", "news" ]
}
```

Collection

Using NoSQL (MongoDB)

- To display the database you are using, type db

```
db
```

```
> db  
test
```

Using NoSQL (MongoDB)

- To list the available databases, type 'show dbs'

```
show dbs
```

```
> show dbs
admin          0.000GB
local          0.000GB
myNewDatabase  0.000GB
```

Using NoSQL (MongoDB)

- To switch databases, issue the use command
- If the database does not exist, it will be created

```
use mymongodatabase
```

```
> use mymongodatabase  
switched to db mymongodatabase
```

Using NoSQL (MongoDB)

- To see the list of collections in the current database, use the 'show collections' command:

```
show collections
```

```
> use local  
switched to db local  
> show collections  
startup_log
```



MONGODB CRUD OPERATIONS



Using NoSQL (MongoDB)

- Create or insert operations, add new documents to a collection.
- If the collection does not currently exist, insert operations will create the collection.
- MongoDB provides the following methods to insert documents into a collection.
 - **db.collection.insertOne()**
 - **db.collection.insertMany()**

Using NoSQL (MongoDB)

```
db.users.insertOne(  ← collection
{
  name: "sue",        ← field: value
  age: 26,            ← field: value
  status: "pending"   ← field: value
}                    } document
)
```

```
> db.users.insertOne( {name: "sue",age: 26,status: "pending"})
{
  "acknowledged" : true,
  "insertedId" : ObjectId("5a02d3541c1652812ae9874c")
}
```

Using NoSQL (MongoDB)

```
db.users.insertOne({name: "sue",age: 26,status: "pending"})
```

```
> db.users.insertMany([{name: "sue",age: 26,status: "pending"},{name: "zan",age: 18,status: "pending"},{name: "don",age: 20,status: "pending"}])
{
  "acknowledged" : true,
  "insertedIds" : [
    ObjectId("5a03c2e11c1652812ae9874f"),
    ObjectId("5a03c2e11c1652812ae98750"),
    ObjectId("5a03c2e11c1652812ae98751")
  ]
}
```


Using NoSQL (MongoDB)

```
db.users.insertMany([{"name": "sue",age: 26,status: "pending"}, {"name": "zan",age: 18,status: "pending"}, {"name": "don",age: 20,status: "pending"}])
```

```
> db.users.insertMany([{"name": "sue",age: 26,status: "pending"}, {"name": "zan",age: 18,status: "pending"}, {"name": "don",age: 20,status: "pending"}])
{
  "acknowledged" : true,
  "insertedIds" : [
    ObjectId("5a03c2e11c1652812ae9874f"),
    ObjectId("5a03c2e11c1652812ae98750"),
    ObjectId("5a03c2e11c1652812ae98751")
  ]
}
```

Using NoSQL (MongoDB)

- Read operations retrieves documents from a collection; i.e. queries a collection for documents.
- MongoDB provides the following methods to read documents from a collection
 - **db.collection.find()**
- You can specify query filters or criteria that identify the documents to return

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
) .limit(5)
```

← collection
← query criteria
← projection
← cursor modifier

Using NoSQL (MongoDB)

```
db.user.find({gender: "female"}, {name: 1}).limit(5)
```

```
> db.user.find({gender: "female"}, {name: 1}).limit(5)
{ "_id" : ObjectId("5b42144ad2963edac48e1997"), "name" : "Dora" }
{ "_id" : ObjectId("5b42154bd2963edac48e1999"), "name" : "Mary" }
```

```
db.users.find({age: {$gt: 18}}, {name: 1, address: 1}).limit(5)
```

```
> db.users.find({age: {$gt: 18}}, {name: 1, address: 1}).limit(5)
{ "_id" : ObjectId("5a02d3541c1652812ae9874c"), "name" : "sue" }
{ "_id" : ObjectId("5a03bf841c1652812ae9874e"), "name" : "ann" }
```

Using NoSQL (MongoDB)

- Update operations modify existing documents in a collection.
- MongoDB provides the following methods to update documents of a collection
 - **db.collection.updateOne()**
 - **db.collection.updateMany()**
 - **db.collection.replaceOne()**
- In MongoDB, update operations target a single collection.
- All write operations in MongoDB are atomic on the level of a single document.
- You can specify criteria, or filters, that identify the documents to update.
- These filters use the same syntax as read operations.

Using NoSQL (MongoDB)

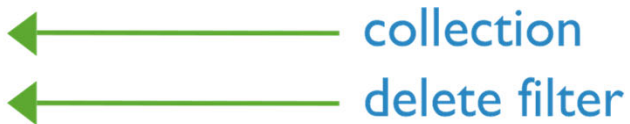
```
db.users.updateMany(  
  { age: { $lt: 18 } },  
  { $set: { status: "reject" } }  
)
```

← collection
← update filter
← update action

Using NoSQL (MongoDB)

- Delete operations remove documents from a collection.
- MongoDB provides the following methods to delete documents of a collection:
 - **db.collection.deleteOne()**
 - **db.collection.deleteMany()**
- In MongoDB, delete operations target a single collection.
- All write operations in MongoDB are atomic on the level of a single document.
- You can specify criteria, or filters, that identify the documents to remove
- These filters use the same syntax as read operations.

```
db.users.deleteMany(  
  { status: "reject" }  
)
```



collection

delete filter

Using NoSQL (MongoDB)

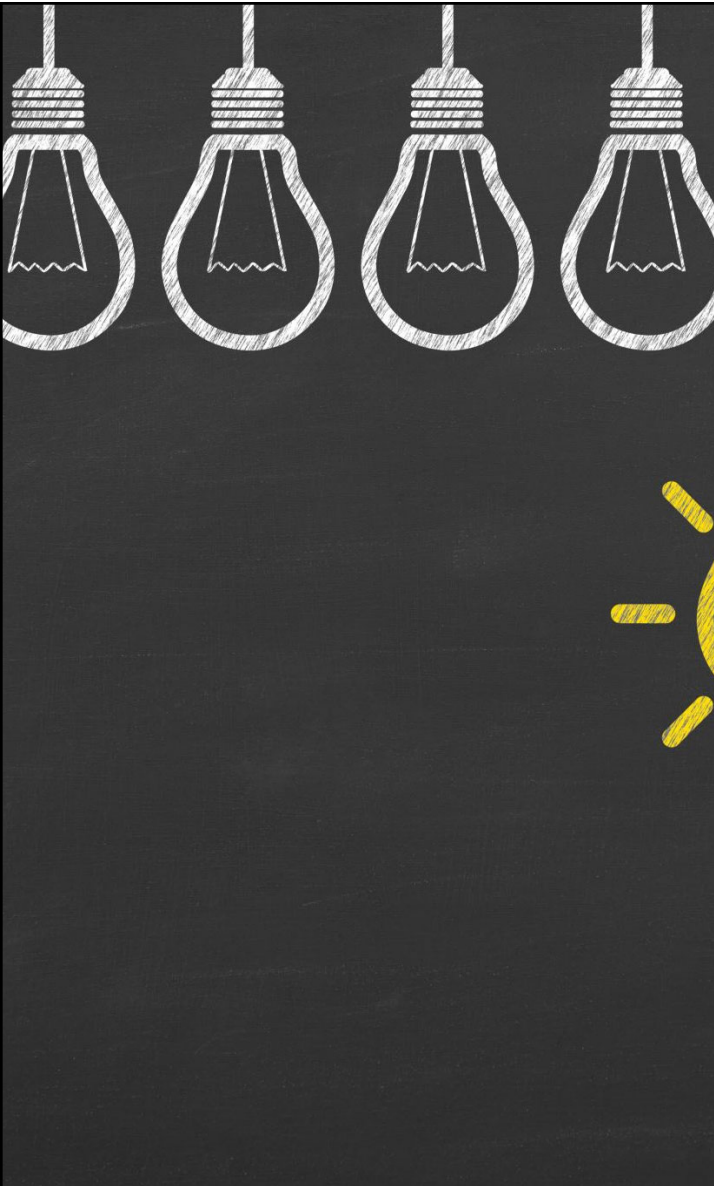
```
db.users.deleteOne({name:"sue"})
```

```
> db.users.deleteOne({name:"sue"})  
{ "acknowledged" : true, "deletedCount" : 1 }
```



NOSQL DATA MODELING





MONGODB DATA MODELING

- MongoDB's document need not have the same schema.
- The document structure can be updated, such as add new fields, remove existing fields, or change the field values to a new type.
- Each field can be a complex structure.
- Data models for MongoDB applications revolves around the structure of documents.
 - Embedded Data Model (Denormalized)
 - Reference Data Model (Normalized)

MONGODB DATA MODELLING : EMBEDDED DATA MODEL

- Store related piece of information in the same document, result in less queries.
- Appropriate:
 - we have “contains” relationships between entities.
 - we have one-to-many relationships between entities. In these relationships the “many” or child documents always appear with or are viewed in the context of the “one” or parent documents.

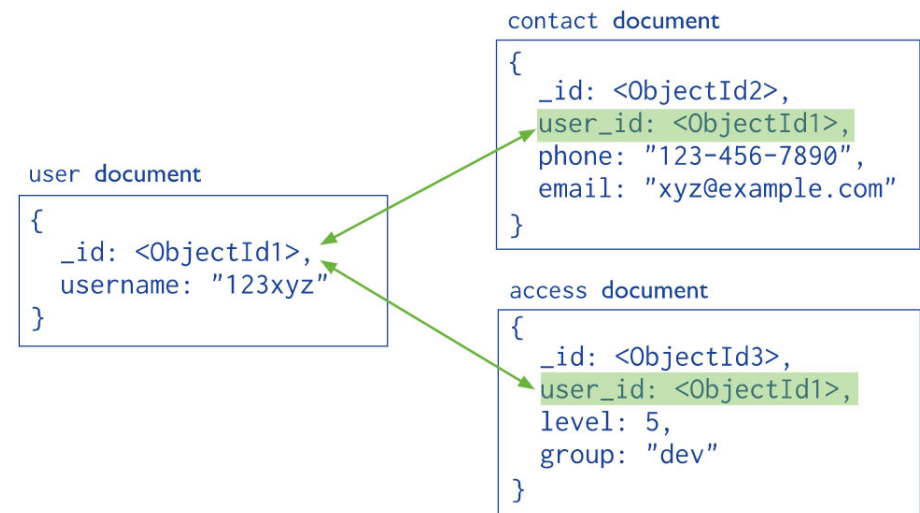
```
{
  _id: <ObjectId>,
  username: "123xyz",
  contact: {
    phone: "123-456-7890",
    email: "xyz@example.com"
  },
  access: {
    level: 5,
    group: "dev"
  }
}
```

Embedded sub-document

Embedded sub-document

MONGODB DATA MODELLING : REFERENCE DATA MODEL

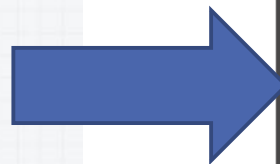
- Store the relationships between data by including links or references from one document to another.
- Appropriate:
 - ❑ when embedding would result in duplication of data but would not provide read performance advantages to outweigh the implications of the duplication.
 - ❑ to represent more complex many-to-many relationships.
 - ❑ to model large hierarchical data sets.



CASE: ONE-TO-ONE MAPPING (EMBEDED)

```
// patron document
{
  _id: "joe",
  name: "Joe Bookreader"
}
```

```
// address document
{
  patron_id: "joe", // reference to patron document
  street: "123 Fake Street",
  city: "Faketon",
  state: "MA",
  zip: "12345"
}
```



```
{
  _id: "joe",
  name: "Joe Bookreader",
  address: {
    street: "123 Fake Street",
    city: "Faketon",
    state: "MA",
    zip: "12345"
  }
}
```

CASE: ONE-TO-MANY MAPPING (EMBEDED)

```
{
  _id: "joe",
  name: "Joe Bookreader"
}

{
  patron_id: "joe",
  street: "123 Fake Street",
  city: "Faketon",
  state: "MA",
  zip: "12345"
}

{
  patron_id: "joe",
  street: "1 Some Other Street",
  city: "Boston",
  state: "MA",
  zip: "12345"
}
```



```
{
  _id: "joe",
  name: "Joe Bookreader",
  addresses: [
    {
      street: "123 Fake Street",
      city: "Faketon",
      state: "MA",
      zip: "12345"
    },
    {
      street: "1 Some Other Street",
      city: "Boston",
      state: "MA",
      zip: "12345"
    }
  ]
}
```

CASE: ONE-TO-MANY

copy

```
{
  title: "MongoDB: The Definitive Guide",
  author: [ "Kristina Chodorow", "Mike Dirolf" ],
  published_date: ISODate("2010-09-24"),
  pages: 216,
  language: "English",
  publisher: {
    name: "O'Reilly Media",
    founded: 1980,
    location: "CA"
  }
}

{
  title: "50 Tips and Tricks for MongoDB Developer",
  author: "Kristina Chodorow",
  published_date: ISODate("2011-05-06"),
  pages: 68,
  language: "English",
  publisher: {
    name: "O'Reilly Media",
    founded: 1980,
    location: "CA"
  }
}
```

CASE: ONE-TO-MANY

PROPOSED SOLUTION A

```
{
  name: "O'Reilly Media",
  founded: 1980,
  location: "CA",
  books: [123456789, 234567890, ...]
}

{
  _id: 123456789,
  title: "MongoDB: The Definitive Guide",
  author: [ "Kristina Chodorow", "Mike Dirolf" ],
  published_date: ISODate("2010-09-24"),
  pages: 216,
  language: "English"
}

{
  _id: 234567890,
  title: "50 Tips and Tricks for MongoDB Developer",
  author: "Kristina Chodorow",
  published_date: ISODate("2011-05-06"),
  pages: 68,
  language: "English"
}
```

[copy](#)

CASE: ONE-TO-MANY

PROPOSED SOLUTION B

```
{
  _id: "oreilly",
  name: "O'Reilly Media",
  founded: 1980,
  location: "CA"
}

{
  _id: 123456789,
  title: "MongoDB: The Definitive Guide",
  author: [ "Kristina Chodorow", "Mike Dirolf" ],
  published_date: ISODate("2010-09-24"),
  pages: 216,
  language: "English",
  publisher_id: "oreilly"
}

{
  _id: 234567890,
  title: "50 Tips and Tricks for MongoDB Developer",
  author: "Kristina Chodorow",
  published_date: ISODate("2011-05-06"),
  pages: 68,
  language: "English",
  publisher_id: "oreilly"
}
```

[copy](#)

MONGODB DATA MODELING: EMBEDDED VS REFERENCE

- Data that is reference together should be embedded.
- Dependent entities can be embedded.
- Entities with one-to-one relationship can be embedded.
- Entities updated at the same time can be embedded.
- Independent entities should be reference.

THE END